

Contemporary Islam

QUANTITATIVE TEXT ANALYSIS REPORT

DATE	March 14, 2026
ANALYSIS ID	019cec92
SCOPE	Model: 2 groups
FEATURES	0 analytics applied

100	179	146.5K	0.73	19
DOCUMENTS	TOTAL CHUNKS	TOTAL WORDS	AVG. SENTIMENT	TOPICS

REPORT NOTICE

This report contains AI-generated text and analytical commentary. AI outputs may contain inaccuracies, fabricated references, or unintended bias. Content is provided for research and informational purposes only and should be independently verified before reliance.

AI-GENERATED SECTIONS

This report combines deterministic statistical analysis with narrative text generated by large language models (LLMs). The table below discloses every AI-generated section included in this report and the model that produced it. AI outputs may contain inaccuracies, fabricated references, or unintended bias and should be independently verified before reliance. All other sections are computed deterministically from the analyzed corpus.

AI-GENERATED SECTION	LOCATION IN REPORT	GENERATED BY
Executive Summary	Executive Summary section	anthropic/claude-opus-4.6 (openrouter)
Composite Cross-Group Comparison	Composite Cross-Group Comparison section	anthropic/claude-opus-4.6 (openrouter)
Conclusion & Recommendations	Conclusion and Recommendations section	anthropic/claude-opus-4.6 (openrouter)

TABLE OF CONTENTS

- 1. Composite Cross-Group Comparison 4
- 2. Executive Summary 7
- 3. Conclusion and Recommendations 11

1. COMPOSITE CROSS-GROUP COMPARISON

Comprehensive Cross-Group Comparison: gemini-3-pro-preview vs. gpt-5.2

1. Executive Overview

This analysis compares content generated by **gemini-3-pro-preview** and **gpt-5.2** across the Contemporary Islam project, spanning word frequency, TF-IDF, sentiment, topic modeling, n-grams, co-occurrence, NER, text classification, network analysis, chunking, similarity, and framing/bias detection. The overarching finding is one of **remarkable structural convergence with meaningful stylistic and distributional differences**. Both providers cover the same thematic terrain -- Islamic finance, human rights, modernity, feminism, fashion, and environmentalism -- but diverge in vocabulary emphasis, writing style, topical granularity, and rhetorical strategy. **gemini-3-pro-preview produces more voluminous, rhetorically assertive text**, while **gpt-5.2 generates more concise, structurally organized, and community-oriented content**.

2. Key Cross-Group Differences

Vocabulary volume and emphasis. gemini-3-pro-preview's top term "islamic" appears 627 times versus gpt-5.2's 402 -- a **56% higher frequency**. gemini-3-pro-preview also uses "muslim" nearly three times as often (161 vs. 57). Conversely, gpt-5.2 emphasizes "religious" significantly more (251 vs. 175, a **43% increase**), and introduces distinctive high-frequency terms absent from gemini-3-pro-preview's top 20: "legal" (134), "community" (100), "modernity" (96), "public" (78), "social" (73), and "avoid" (60). This suggests **gpt-5.2 favors broader sociological and practical framing**, while **gemini-3-pro-preview defaults more heavily to identity-label terminology**.

TF-IDF distinctiveness. While both share "halal" as the top TF-IDF term (~4.8), gpt-5.2 gives "circular" a substantially higher TF-IDF weight (3.78 vs. 2.65) and elevates "modernity" (2.50 vs. 1.37), "human rights" (2.21 vs. 1.30), "community" (1.37 vs. absent), and "online" (1.25 vs. absent). gemini-3-pro-preview uniquely surfaces "fashion" (2.60), "economy" (2.09), "shura" (1.77), "green" (1.64), and "extremist" (1.29). **gpt-5.2's TF-IDF profile is more legally and conceptually oriented; gemini-3-pro-preview's is more thematically concrete**.

Writing style and framing. The framing analysis reveals the starkest stylistic divergence. gemini-3-pro-preview employs **2.5× more passive voice** (199 vs. 79), **6× more intensifiers** (31 vs. 5), and **28% more loaded terms** (269 vs. 210). gpt-5.2 uses slightly more hedging (17 vs. 15). Complexity grades are virtually identical (~15.1). This pattern indicates **gemini-3-pro-preview adopts a more assertive, declarative academic tone**, while **gpt-5.2 writes with greater restraint and hedged precision**.

Text classification. gemini-3-pro-preview classifies 29.3% of content as Economy & Business versus gpt-5.2's 17.5%. gpt-5.2 assigns a far larger share to **Law & Security** (32.5% vs. 20.2%) and **Politics** (10.0% vs. 5.1%). gpt-5.2 also spreads across more categories (13 vs. 10), including Education, Health, and Quran & Revelation -- absent from gemini-3-pro-preview's output.

NER entity distribution. gemini-3-pro-preview identifies far more entities overall, particularly NORP (1,339 vs. 801), PERSON (426 vs. 32), GPE (292 vs. 25), and LAW (79 vs. 9). gpt-5.2's MONEY entities spike (523 vs. 423), driven by structured numbered formatting (### 1, ### 2, etc.) misclassified as monetary entities -- a notable **artifact of gpt-5.2's list-heavy formatting style**.

3. Cross-Feature Patterns

Consistent pattern: gpt-5.2 emphasizes legal/community framing. The word "legal" (134 in gpt-5.2, absent from gemini-3-pro-preview's top 20), "community" (100 vs. absent), and "public" (78 vs. absent) align with gpt-5.2's higher TF-IDF for "human rights" and its classification lean toward Law & Security. Topic modeling reinforces this: gpt-5.2 produces a dedicated "Human Rights and Islamic Law" topic at 12.6% prevalence with "legal" as its third keyword.

Consistent pattern: gemini-3-pro-preview is more geographically and personally specific. NER shows gemini-3-pro-preview identifying 292 GPE entities (Indonesia 31, Turkey 15, Europe 14) versus gpt-5.2's 25, and 426 PERSON entities (Muhammad 26, Maqasid al-Sharia 14) versus 32. This aligns with gemini-3-pro-preview's higher "muslim" frequency and suggests **gemini-3-pro-preview grounds arguments in specific people, places, and historical contexts**, while **gpt-5.2 abstracts toward principles and frameworks**.

Reinforcing finding: Topic granularity. gpt-5.2's LDA generates 13 interpretable topics versus gemini-3-pro-preview's 6 dominant ones. gpt-5.2 separates "Islam and Feminism Discourse" (5.0%) from "Human Rights and Islamic Law" (12.6%), while gemini-3-pro-preview merges these into a broader "Islam, Modernity, and Social Reform" (37.4%). Both providers' dominant macro-theme clusters (~70-96%) center on Islamic/religious/modern content, but **gpt-5.2 provides finer-grained topical differentiation**.

Contradiction: Sentiment vs. Framing. Both providers are overwhelmingly positive in sentiment (gemini-3-pro-preview 86.9%, gpt-5.2 88.8%), yet both show substantial loaded terms (269/210). This tension suggests positivity manifests through **value-laden advocacy language** rather than neutral reporting.

4. Similarities & Convergences

- **Co-occurrence networks are identical:** same 50 nodes, 728 edges, density 0.5943, clustering 0.7080. The underlying source corpus drives identical structural relationships.
- **Top co-occurrence pairs match exactly** (ijarah-leasing, artificial-intelligence, street-style, etc.), confirming shared source material.
- **Top n-grams converge:** "references external sources" and "external sources used" lead both groups, revealing shared boilerplate reference sections.
- **Similarity clusters parallel:** both form halal (7-8 docs), circular economy (7-8 docs), human rights (7 docs), and fashion (5-6 docs) clusters at comparable internal similarity scores (~0.20-0.30).
- **Sentiment polarity** is nearly identical, with negligible variation (<2.5 percentage points).

5. Strategic Implications

- **Content differentiation is stylistic, not substantive.** Both providers cover identical topics from the same sources. Stakeholders seeking **varied topical coverage** should not expect meaningful differences between providers.
- **gemini-3-pro-preview is preferable for richly contextualized, historically grounded content** -- it names more people, places, and specific legal frameworks, and clusters content into broader narrative arcs.
- **gpt-5.2 is preferable for structured, legally precise, community-oriented content** -- its list-based formatting, legal vocabulary emphasis, and finer topic segmentation suit policy analysis and practical guidance documents.

- **Bias monitoring should focus on gemini-3-pro-preview**, which uses 2.5× more passive constructions and 6× more intensifiers, potentially obscuring agency or amplifying claims.
- **Both providers generate near-duplicate boilerplate** ("No external sources used") that inflates similarity scores and should be filtered in production workflows.
- **gpt-5.2's formatting artifacts** (numbered headers misclassified as MONEY entities) require post-processing cleanup for NER-dependent applications.

2. EXECUTIVE SUMMARY

2.1. gemini-3-pro-preview Analysis

gemini-3-pro-preview's Narrative Construction of Contemporary Islam: A Computational Synthesis

This analysis examines 50 AI-generated documents (99 chunks, ~79,100 words) produced by Google's gemini-3-pro-preview model on the theme of contemporary Islam. The corpus spans topics including radicalism, feminism, halal standards, Islamic banking, human rights, technology, environment, fashion, modernity, and democracy. Through twelve computational features, we can characterize how gemini-3-pro-preview frames Islam's relationship to modern global issues, identify its dominant rhetorical strategies, and surface both its strengths and blind spots as a knowledge synthesizer.

Macro-Theme Architecture

Consolidating topic modeling, text classification, and chunking results, gemini-3-pro-preview's output organizes into four primary macro-clusters:

1. **Islamic Economics & Sustainability** (dominant, ~33% of topic weight + 29.3% classification): Circular economy, Islamic finance, green sukuk, and ethical economics form the largest thematic cluster. gemini-3-pro-preview treats Islam's economic principles as inherently compatible with sustainability discourse.
2. **Islam, Modernity & Social Reform** (~37.4% topic weight): The single largest topic by weight, blending women's rights, Quranic interpretation, secularism, and modernity. This cluster absorbs feminism, gender advocacy, and theological modernization.
3. **Law, Rights & Security** (~20.5% topic weight + 20.2% classification): Human rights, the Cairo Declaration, deradicalization, and extremism. This is the most negatively sentimentalized cluster.
4. **Identity, Practice & Material Culture** (~9% topic weight): Halal certification, modest fashion, digital religious practice, and lifestyle. This is the thinnest cluster, suggesting gemini-3-pro-preview underweights lived religious experience relative to institutional discourse.

Cross-Feature Synthesis

Several convergent findings emerge when we read across features:

gemini-3-pro-preview's signature move is harmonization. The overwhelming positive sentiment (86.9%), combined with high loaded-term density (269 instances) and low hedging (15 instances), reveals a consistent rhetorical pattern: gemini-3-pro-preview frames Islam as fundamentally compatible with modern values while using assertive, often value-laden language. This is not neutral academic synthesis; it is advocacy-adjacent framing dressed in academic register. The average text complexity grade of 15.07 (graduate-level) reinforces this impression of authoritative, but not necessarily balanced, discourse.

The co-occurrence network reveals a concept vocabulary rooted in Islamic jurisprudence. Pairs like *ijarah*-leasing, *consensus-ijma*, *ijtihad-reasoning*, *balance-mizan*, and *al-sharia-maqasid* show that gemini-3-pro-preview systematically introduces Arabic theological and legal terminology alongside English glosses. This is a distinctive feature: gemini-3-pro-preview appears to function as a bilingual glossary embedded within its narratives, which could be pedagogically valuable but also creates an impression of depth that may exceed the actual analytical rigor.

Redundancy and template effects are significant. The n-gram analysis surfaces meta-textual phrases like "references external sources," "external sources used," and "references external" across multiple documents.

These are not content phrases but gemini-3-pro-preview's self-referential commentary on its own sourcing practices, appearing as a formulaic disclaimer. This inflates apparent phrase diversity without adding substantive content and should be treated as a methodological artifact.

The text network is dense (0.594) and fully connected (1 component, 50 nodes, 728 edges) with high clustering (0.708). This means gemini-3-pro-preview's documents are thematically interwoven rather than siloed. No document is conceptually isolated. While this suggests coherence, it may also reflect homogeneity: gemini-3-pro-preview applies a similar vocabulary and framing strategy across all topics rather than adapting its register to domain-specific requirements.

Research Gaps and Silent Voices

Several significant absences emerge:

- **Sectarian diversity is invisible.** There is no meaningful differentiation between Sunni, Shia, Sufi, Ibadi, or other traditions. gemini-3-pro-preview constructs a monolithic "Islam" that erases internal theological pluralism.
- **Geographic specificity is thin.** While "Middle East" appears as a co-occurrence pair, there is little engagement with specific national contexts, diaspora experiences, or regional variations in Islamic practice.
- **Critical and dissenting Muslim voices are absent.** The corpus lacks representation of ex-Muslim perspectives, intra-Muslim critiques, or heterodox positions. The harmonization bias suppresses the very debates that characterize actual Islamic intellectual life.
- **Worship and ritual practice receive minimal attention** (1.0% classification), which is striking for a corpus ostensibly about Islam. gemini-3-pro-preview prioritizes Islam as a system of ethics and governance over Islam as a lived spiritual practice.

Methodological Caveats

- All 50 documents are AI-generated, meaning we are analyzing gemini-3-pro-preview's model of Islamic discourse, not Islamic discourse itself. Patterns reflect training data biases and alignment tuning.
- The high positive sentiment is likely an artifact of gemini-3-pro-preview's alignment toward constructive, non-offensive output rather than a genuine finding about Islamic scholarship.
- Topic modeling with overlapping keyword sets (e.g., "islamic" and "environmental" appearing across multiple topics) suggests the corpus may have been under-differentiated for optimal LDA performance.
- Co-occurrence and n-gram NPMI scores for low-frequency items (e.g., "references external sources" with frequency 1) should be interpreted with extreme caution, as statistical reliability requires higher counts.

Actionable Recommendations

1. **For comparative analysis across AI providers:** Use the harmonization bias, loaded-term density, and Arabic terminology integration as benchmarking variables. These are gemini-3-pro-preview's most distinctive features and will likely diverge from other providers.
2. **For literature review synthesis:** Treat gemini-3-pro-preview's output as a starting framework for identifying topics but supplement with actual scholarly sources, particularly for contested areas like radicalism, human rights, and feminism where gemini-3-pro-preview's framing suppresses debate.
3. **For bias auditing:** The gap between high loaded-term counts (269) and low hedging (15) is the most actionable signal. gemini-3-pro-preview presents contentious claims with high confidence, which could mislead readers about the state of scholarly consensus.

2.2. gpt-5.2 Analysis

Corpus Overview: gpt-5.2's Narrative Construction of Contemporary Islam

This analysis examines 50 documents (80 chunks, ~67,400 words) generated by gpt-5.2 on topics spanning contemporary Islam. The corpus was designed to probe how this AI provider frames Islam in relation to modern global issues including radicalism, feminism, halal economy, Islamic banking, human rights, technology, environment, fashion, modernity, and democracy.

Executive Synthesis

OpenAI's output on contemporary Islam is characterized by four overarching findings:

1. A Legal-Security Frame Dominates the Corpus. Text classification reveals that 32.5% of all chunks are categorized under "Law & Security," making it the single largest category--nearly double the next-largest, "Economy & Business" (17.5%). This signals that gpt-5.2's default framing of Islam-related content gravitates toward juridical and security-oriented language, even when the underlying topics (fashion, environment, halal food) are not inherently about law or security. This is a significant framing effect: it suggests the model has internalized a discourse in which Islam is primarily discussed through regulatory, legal, and rights-based lenses.

2. Overwhelmingly Positive Sentiment Masks Complexity. Nearly 89% of chunks are classified as positive, with only 11.3% negative and zero neutral. For a corpus addressing radicalism, human rights tensions, and misconceptions, this distribution is analytically suspicious. It likely reflects gpt-5.2's alignment tuning--a tendency to produce balanced, conciliatory, and constructive prose rather than critically negative or genuinely neutral analysis. The absence of any neutral text is a strong artifact of model behavior, not a reflection of the source material's intellectual character.

3. Thematic Breadth with Conceptual Redundancy. Topic modeling surfaces 10 labeled topics, but several overlap substantially. For instance, "Islamic Environmentalism and Green Living" appears twice with different keyword mixes (one emphasizing climate/water, the other circular economy/sukuk). Similarly, "Halal Economy and Online Islam" and "Halal Food and Lifestyle Markets" share significant vocabulary. This redundancy--confirmed by the high network density (0.594) and clustering coefficient (0.708)--indicates that gpt-5.2 tends to recycle a common vocabulary across nominally distinct topics, reducing genuine thematic differentiation.

4. Arabic/Islamic Terminology Is Present but Instrumentally Deployed. Co-occurrence analysis surfaces terms like *ijarah-leasing*, *ijma-consensus*, *ijtihad-reasoning*, *mizan-balance*, *maqasid al-sharia*, and the *Cairo Declaration*. These terms always appear paired with their English glosses, suggesting gpt-5.2 uses Arabic terminology as explanatory devices rather than engaging with them as contested intellectual concepts. This is pedagogically useful but analytically shallow.

Macro-Theme Clusters

Consolidating the 10 topic-model outputs and cross-referencing with classification, n-grams, and chunking, six higher-order clusters emerge:

- **Macro-Theme 1: Law, Rights, and Governance** (topics on human rights, Islamic law, democracy, liberal democracy) -- the most prominent cluster, accounting for roughly 30-35% of corpus weight.
- **Macro-Theme 2: Halal Economy and Ethical Markets** (halal food, halal certification, Islamic finance, circular economy) -- the second-largest cluster (~25%).

- **Macro-Theme 3: Environment and Stewardship** (Islamic environmentalism, climate action, zakat/waqf for climate) -- a distinct but internally repetitive cluster (~10%).
- **Macro-Theme 4: Gender, Feminism, and Modesty** (women's rights, feminist discourse, modest fashion) -- moderately represented but with high bias scores.
- **Macro-Theme 5: Technology and Digital Islam** (responsible tech, online radicalization, digital halal products) -- present but thin.
- **Macro-Theme 6: Modernity and Identity** (Islamic modernity, faith/secular reconciliation, misconceptions) -- a cross-cutting frame rather than a standalone cluster.

Research Gaps and Silent Voices

Several important dimensions of contemporary Islam are absent or marginalized in the gpt-5.2 corpus:

- **Sectarian diversity** (Sunni/Shia/Sufi distinctions) receives no visible treatment.
- **Lived religious experience** (prayer, fasting, pilgrimage as phenomenological rather than legal categories) is nearly invisible--"Worship & Ritual Practice" accounts for only 3.8%.
- **Muslim-majority country specificity** beyond generic references is minimal; the named entity profile is thin on specific political figures, scholars, or institutions.
- **Critical and dissenting Muslim voices** (progressive Islam, ex-Muslim perspectives, internal reform movements) appear absent.
- **Colonialism and postcolonial context** as a frame for understanding contemporary Islam is not surfaced in any topic or n-gram cluster.

Methodological Caveats

- The corpus is entirely AI-generated, meaning all findings describe gpt-5.2's narrative construction, not empirical reality. Frequency patterns reflect model priors, not scholarly consensus.
- The high redundancy across topics may partly be a modeling artifact (LDA struggling with a relatively small, homogeneous corpus), but it also reflects genuine repetitiveness in gpt-5.2's prose.
- Sentiment analysis on AI-generated text is circular: the model produces text designed to be balanced, and then we measure that balance. The 89% positive finding should be read as a measure of alignment tuning, not content quality.
- Network metrics showing zero values for top nodes by degree and betweenness suggest a data reporting issue or an extremely uniform distribution of centrality, consistent with a corpus where no single concept dominates all others but many share overlapping vocabulary.

3. CONCLUSION AND RECOMMENDATIONS

Conclusion

The comparative analysis of gemini-3-pro-preview and gpt-5.2 outputs across the Contemporary Islam project reveals a corpus that is thematically unified but structurally differentiated by provider. Both groups engage the same source material and converge on core thematic pillars -- Islamic finance and the circular economy, human rights under Islamic law, modest fashion, feminism, political governance, environmental stewardship, and the halal economy. However, meaningful differences emerge in vocabulary distribution, topical granularity, classification emphasis, entity recognition depth, content redundancy patterns, and rhetorical framing.

Vocabulary and Term Weighting

gemini-3-pro-preview produces a notably higher raw frequency for the anchor term "islamic" (627 vs. 402), suggesting longer or more repetitive elaboration around core Islamic concepts. gpt-5.2, by contrast, surfaces terms absent from gemini-3-pro-preview's top-20 -- "legal" (134), "community" (100), "public" (78), "social" (73), and "avoid" (60) -- pointing to a more practical, action-oriented, and community-centered lexicon. TF-IDF analysis reinforces this: gemini-3-pro-preview weights "islamic" (3.42) far above gpt-5.2 (1.81), while gpt-5.2 elevates "circular" (3.78 vs. 2.65), "modernity" (2.50 vs. 1.37), "human rights" (2.21 vs. 1.30), and "community" (1.37, absent in gemini-3-pro-preview's top-20). gpt-5.2 thus distributes semantic emphasis more evenly across sub-themes, whereas gemini-3-pro-preview concentrates weight on the overarching "Islamic" identifier.

Topic Structure and Coverage

gemini-3-pro-preview's LDA model yields six primary topics with a single macro-cluster absorbing 95.8% of prevalence (Islam/Modernity/Social Reform at 37.4%), leaving democracy and political pluralism as an outlier at only 4.2%. gpt-5.2 resolves into 13 topics with more balanced distribution: Islamic Governance and Political Modernity (26.7%), Halal Finance and Circular Economy (14.5%), Human Rights and Islamic Law (12.6%), and several mid-range topics between 5-10%. This finer-grained decomposition makes gpt-5.2's output more navigable for researchers seeking discrete thematic entry points. Both groups flag overlapping topics requiring consolidation, particularly around religious identity, legal frameworks, and community/digital practice.

Sentiment and Classification

Sentiment distributions are closely aligned -- gemini-3-pro-preview at 86.9% positive and gpt-5.2 at 88.8% -- confirming that both providers maintain a constructive, advocacy-adjacent tone. Classification diverges more sharply: gemini-3-pro-preview assigns 29.3% of content to Economy & Business and 20.2% to Law & Security, while gpt-5.2 inverts the weighting (17.5% Economy & Business, 32.5% Law & Security) and activates additional categories including Education (3.8%), Health (1.3%), and Quran & Revelation (1.3%). gpt-5.2's broader classification schema captures more disciplinary diversity from the same corpus.

Entity Recognition and Structural Signals

gemini-3-pro-preview identifies significantly more named entities overall (4,316 vs. 2,173), with richer geographic (Indonesia 31, Turkey 15, Europe 14) and person-level tagging (Muhammad 26, Maqasid al-Sharia 14). gpt-5.2's NER is comparatively sparse -- only one person entity (Islam, 11) and one geographic entity reached the top lists -- and over-generates MONEY-type entities through heading markers (###), suggesting less robust preprocessing.

Co-occurrence networks are identical between groups (density 0.59, clustering 0.71), confirming shared structural topology at the network level.

Redundancy, N-Grams, and Framing

gpt-5.2 exhibits three near-duplicate chunk pairs (including one perfect 1.0 similarity) versus gemini-3-pro-preview's single near-duplicate, driven by identical "No external sources used" boilerplate. gpt-5.2's n-gram profile is substantially richer (16 significant collocations vs. 6), surfacing domain-specific phrases like "human rights lens," "goes wrong," "street style," and "actionable checks" that indicate more varied phraseological output. On framing, gemini-3-pro-preview uses far more passive voice (199 vs. 79) and intensifiers (31 vs. 5), while gpt-5.2 relies slightly more on hedging (17 vs. 15). Both average a college-level complexity grade (~15), but gemini-3-pro-preview's higher passive and intensifier counts suggest a more formal, sometimes less direct rhetorical style. High-bias resource profiles differ: gemini-3-pro-preview flags theology-and-extremism pieces, while gpt-5.2 flags radicalization and intersectional rights content.

Recommendations

High Priority

Deduplicate boilerplate content across gpt-5.2 outputs. Three near-duplicate pairs -- including an exact 1.0 match -- stem from "No external sources used" reference blocks. These inflate similarity metrics and distort redundancy analysis. Implement post-generation stripping or chunking rules that exclude formulaic reference sections before analysis.

Standardize NER preprocessing for gpt-5.2. The heavy MONEY-type entity counts driven by markdown heading symbols (###, #) indicate that gpt-5.2's raw output is not being adequately cleaned before entity extraction. Apply regex-based header stripping to ensure NER results reflect genuine named entities rather than formatting artifacts.

Leverage gpt-5.2's finer topic resolution for thematic navigation. gpt-5.2's 13-topic model provides more actionable segmentation than gemini-3-pro-preview's 6-topic model, where a single macro-cluster dominates. For research portals or content indexes targeting Contemporary Islam, adopt gpt-5.2's topic structure as the primary navigation scaffold, supplemented by gemini-3-pro-preview's broader contextual framing.

Medium Priority

Reduce passive voice density in gemini-3-pro-preview outputs. At 199 passive constructions compared to gpt-5.2's 79, gemini-3-pro-preview content may feel less direct and harder to parse for general audiences. If these outputs serve educational or public-facing purposes, apply editorial guidelines or post-processing prompts that favor active constructions.

Enrich gpt-5.2's geographic and biographical entity coverage. gemini-3-pro-preview identifies 31 Indonesia references, 15 Turkey, and 26 Muhammad mentions; gpt-5.2 surfaces almost none of these. For projects requiring geopolitical or historical-figure mapping, either supplement gpt-5.2 outputs with gemini-3-pro-preview-derived entity layers or adjust gpt-5.2 prompting to elicit more specific proper nouns.

Consolidate overlapping topics flagged by both providers. Both LDA models identify redundancy between religious identity, legal frameworks, and digital community topics. Merging these into consolidated super-topics

would reduce noise and improve coherence scores, particularly for gemini-3-pro-preview's zero-prevalence outlier topics (Topics 7-10).

Low Priority

Expand gpt-5.2's classification taxonomy for reuse. gpt-5.2 activates 13 classification categories versus gemini-3-pro-preview's 10, including Education, Health, and Quran & Revelation. Consider adopting this broader schema as a project-wide standard to capture disciplinary nuances that gemini-3-pro-preview's coarser classification misses.

Monitor loaded-term density in sensitive sub-corpora. Both providers concentrate high-bias scores in radicalization, extremism, and rights-intersection content. For downstream publication or training use, flag these resources for manual review to ensure balanced framing, particularly gemini-3-pro-preview's "Distinguishing mainstream theology from extremist ideology" (bias score 5.55) and gpt-5.2's "How online networks accelerate Islamic radicalization" (bias score 5.87).